

# Lecture 11:

## CNNs in Practice



# Administrative

- Midterms are graded!
  - Pick up now
  - Or in Andrej, Justin, Albert, or Serena's OH
- Project milestone due today, 2/17 by midnight
  - Turn in to Assignments tab on Coursework!
- Assignment 2 grades soon
- Assignment 3 released, due 2/24

# Midterm stats

**Mean: 75.0    Median: 76.3**  
**N: 311        Max: 103.0**


**Standard Deviation: 13.2**



# Midterm stats




[We threw out TF3 and TF8]




# Midterm stats



Question 3.1 mean scores



Question 3.2 mean scores




# Midterm Stats



**Bonus mean: 0.8**

# Last Time



Recurrent neural networks  
for modeling sequences

## Vanilla RNNs

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

## LSTMs

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g$$

$$h_t^l = o \odot \tanh(c_t^l)$$

# Last Time

PANDARUS:

Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and  
my fair nues begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

**Lemma 0.1.** Assume (3) and (3) by the construction in the description.

Suppose  $X = \lim |X|$  (by the formal open covering  $X$  and a single map  $\text{Proj}_X(A) = \text{Spec}(B)$  over  $U$  compatible with the complex

$$\text{Set}(A) = \Gamma(X, \mathcal{O}_X, \mathcal{O}_X).$$

When in this case of to show that  $\mathcal{Q} \rightarrow \mathcal{C}_{Z/X}$  is stable under the following result in the second conditions of (1), and (3). This finishes the proof. By Definition ?? (without element is when the closed subschemes are catenary. If  $T$  is surjective we may assume that  $T$  is connected with residue fields of  $S$ . Moreover there exists a closed subspace  $Z \subset X$  of  $X$  where  $U$  in  $X'$  is proper (some defining as a closed subset of the uniqueness it suffices to check the fact that the following theorem

(1)  $f$  is locally of finite type. Since  $S = \text{Spec}(R)$  and  $Y = \text{Spec}(R)$ .

*Proof.* This is form all sheaves of sheaves on  $X$ . But given a scheme  $U$  and a surjective étale morphism  $U \rightarrow X$ . Let  $U \cap U = \coprod_{i=1, \dots, n} U_i$  be the scheme  $X$  over  $S$  at the schemes  $X_i \rightarrow X$  and  $U = \lim_i X_i$ .  $\square$

The following lemma surjective restrocomposes of this implies that  $\mathcal{F}_{x_0} = \mathcal{F}_{x_0} = \mathcal{F}_{X, \dots, 0}$ .

**Lemma 0.2.** Let  $X$  be a locally Noetherian scheme over  $S$ ,  $E = \mathcal{F}_{X/S}$ . Set  $\mathcal{I} = \mathcal{I}_1 \subset \mathcal{I}'_n$ . Since  $\mathcal{I}^n \subset \mathcal{I}'^n$  are nonzero over  $i_0 \leq \mathfrak{p}$  is a subset of  $\mathcal{I}_{n,0} \circ \mathcal{A}_2$  works.

**Lemma 0.3.** In Situation ???. Hence we may assume  $\mathfrak{q}' = 0$ .

*Proof.* We will use the property we see that  $\mathfrak{p}$  is the next functor (??). On the other hand, by Lemma ?? we see that

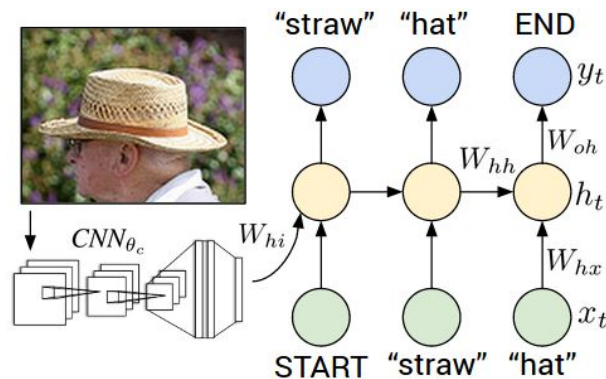
$$D(\mathcal{O}_{X'}) = \mathcal{O}_X(D)$$

where  $K$  is an  $F$ -algebra where  $\delta_{n+1}$  is a scheme over  $S$ .  $\square$

## Sampling from RNN language models to generate text



# Last Time



CNN + RNN for  
image captioning

Cell that robustly activates inside if statements:

```
static int __dequeue_signal(struct sigpending *pending,
                           siginfo_t *info)
{
    int sig = next_signal(pending, mask);
    if (sig) {
        if (current->notifier) {
            if (sigismember(current->notifier_mask, sig)) {
                if (!!(current->notifier)(current->notifier_data)) {
                    clear_thread_flag(TIF_SIGPENDING);
                    return 0;
                }
            }
        }
        collect_signal(sig, pending, info);
    }
    return sig;
}
```

Interpretable RNN cells


# Today

Working with CNNs in practice:

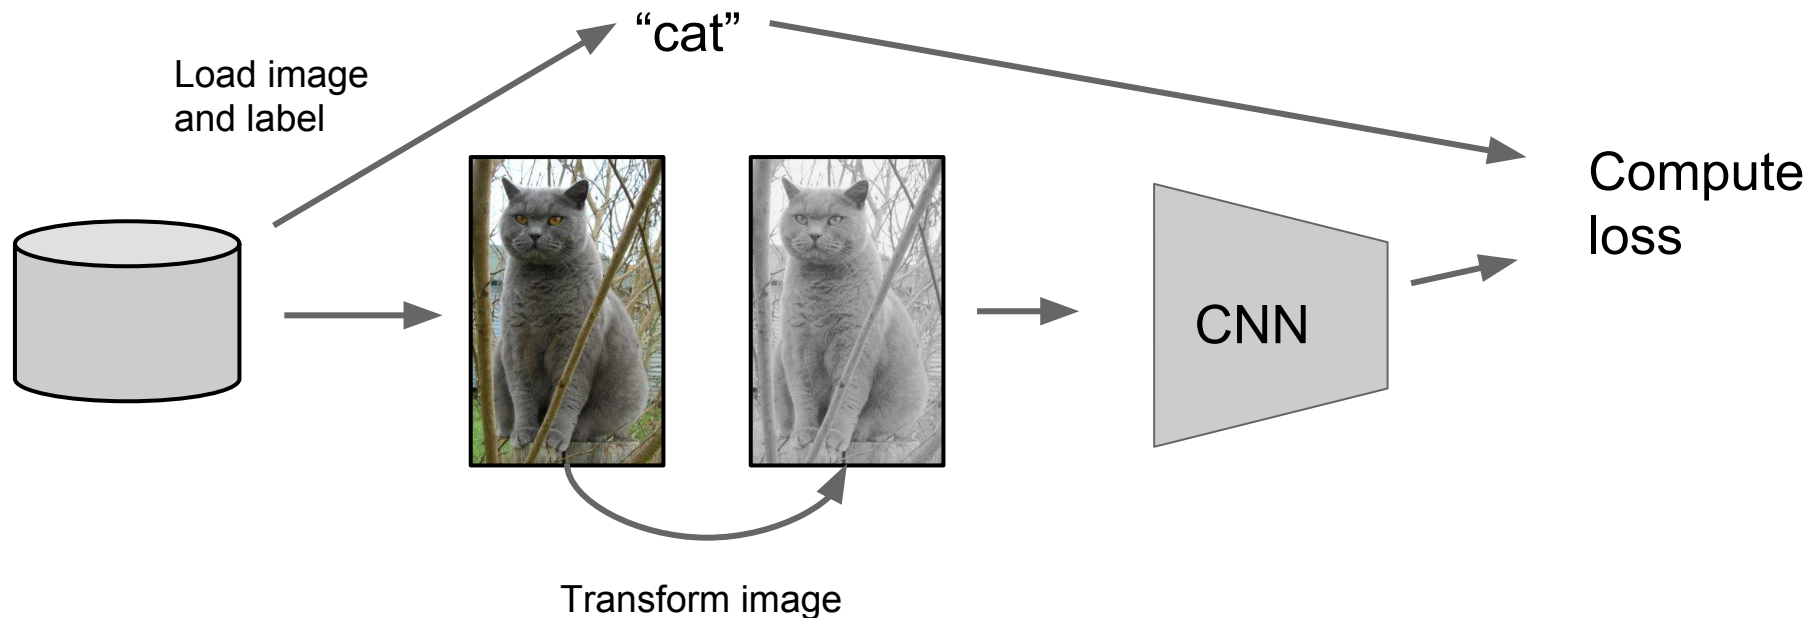
- Making the most of your data
  - Data augmentation
  - Transfer learning
- All about convolutions:
  - How to arrange them
  - How to compute them fast
- “Implementation details”
  - GPU / CPU, bottlenecks, distributed training

# Data Augmentation

# Data Augmentation




# Data Augmentation



# Data Augmentation

- Change the pixels without changing the label
- Train on transformed data
- VERY widely used



What the computer sees

# Data Augmentation


## 1. Horizontal flips



# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales






# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales

ResNet:

1. Pick random  $L$  in range  $[256, 480]$
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch



# Data Augmentation


## 2. Random crops/scales

**Training:** sample random crops / scales

ResNet:

1. Pick random  $L$  in range  $[256, 480]$
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch

**Testing:** average a fixed set of crops



# Data Augmentation

## 2. Random crops/scales

**Training:** sample random crops / scales


ResNet:

1. Pick random  $L$  in range  $[256, 480]$
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch

**Testing:** average a fixed set of crops

ResNet:

1. Resize image at 5 scales:  $\{224, 256, 384, 480, 640\}$
2. For each size, use 10  $224 \times 224$  crops: 4 corners + center, + flips




# Data Augmentation

## 3. Color jitter

**Simple:**

Randomly jitter contrast




# Data Augmentation

## 3. Color jitter

**Simple:**

Randomly jitter contrast



**Complex:**

1. Apply PCA to all [R, G, B] pixels in training set
2. Sample a “color offset” along principal component directions
3. Add offset to all pixels of a training image

(As seen in *[Krizhevsky et al. 2012]*, ResNet, etc)

# Data Augmentation



## 4. Get creative!

Random mix/combinations of :


- translation
- rotation
- stretching
- shearing,
- lens distortions, ... (go crazy)

# A general theme:


1. **Training:** Add random noise
2. **Testing:** Marginalize over the noise



Data Augmentation




(a) Standard Neural Net



(b) After applying dropout.

Dropout



DropConnect

Batch normalization, Model ensembles

# Data Augmentation: Takeaway

- Simple to implement, use it
- Especially useful for small datasets
- Fits into framework of noise / marginalization



# Transfer Learning

“You need a lot of a data if you want to train/use CNNs”

# Transfer Learning

“You need a lot of data if you want to train/use CNNs”

**BUSTED**

# Transfer Learning with CNNs

image

conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

maxpool

FC-4096

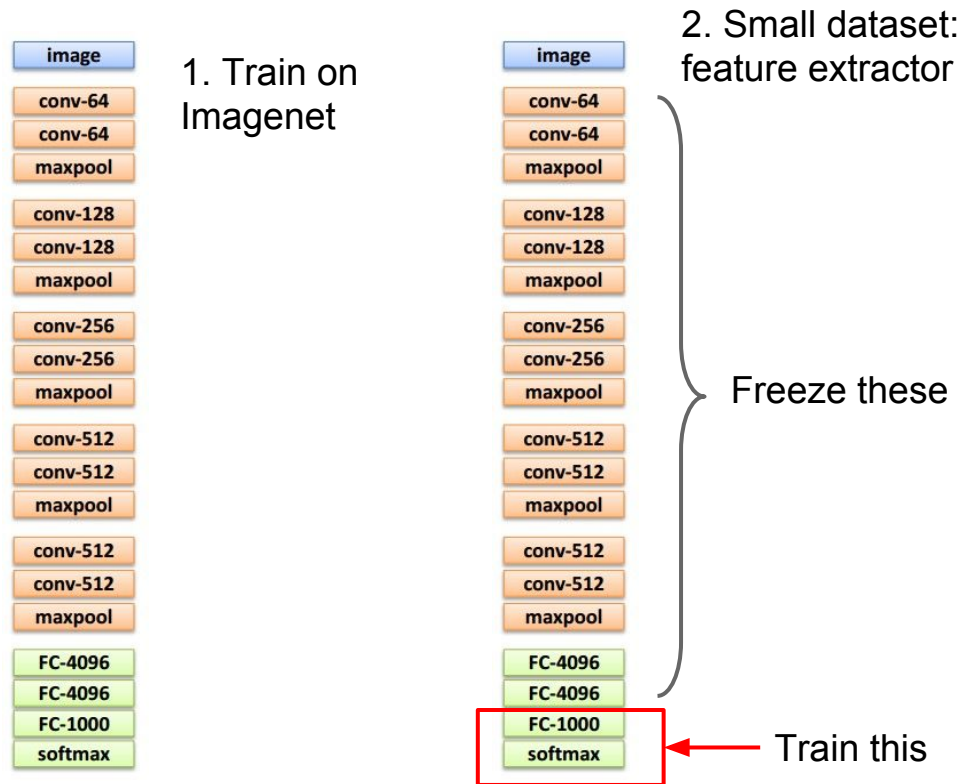
FC-4096

FC-1000


softmax

1. Train on  
Imagenet

# Transfer Learning with CNNs



# Transfer Learning with CNNs




1. Train on  
Imagenet



2. Small dataset:  
**feature extractor**

Freeze these

Train this




3. Medium dataset:  
**finetuning**

more data = retrain more of  
the network (or all of it)


Freeze these

Train this

# Transfer Learning with CNNs




1. Train on Imagenet



2. Small dataset:  
**feature extractor**

Freeze these

Train this



3. Medium dataset:  
**finetuning**

more data = retrain more of the network (or all of it)

Freeze these

tip: use only  $\sim 1/10$ th of the original learning rate in finetuning top layer, and  $\sim 1/100$ th on intermediate layers

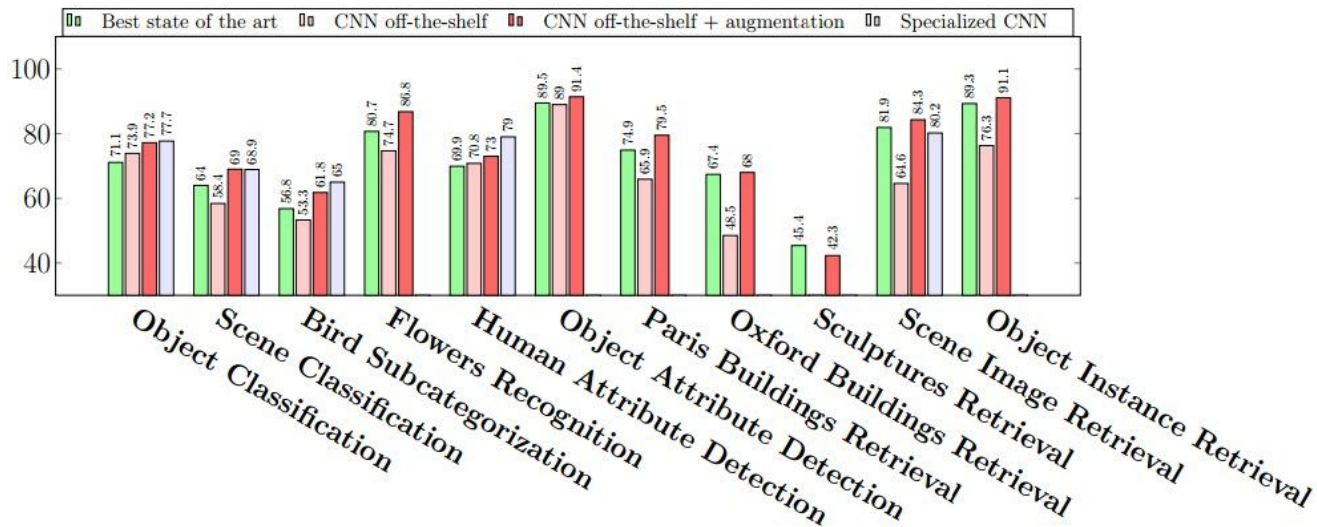

Train this


# CNN Features off-the-shelf: an Astounding Baseline for Recognition

[Razavian et al, 2014]

*DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition*  
 [Donahue\*, Jia\*, et al., 2013]


	DeCAF <sub>6</sub>	DeCAF <sub>7</sub>
LogReg	40.94 ± 0.3	40.84 ± 0.3
SVM	39.36 ± 0.3	40.66 ± 0.3
Xiao et al. (2010)	38.0	



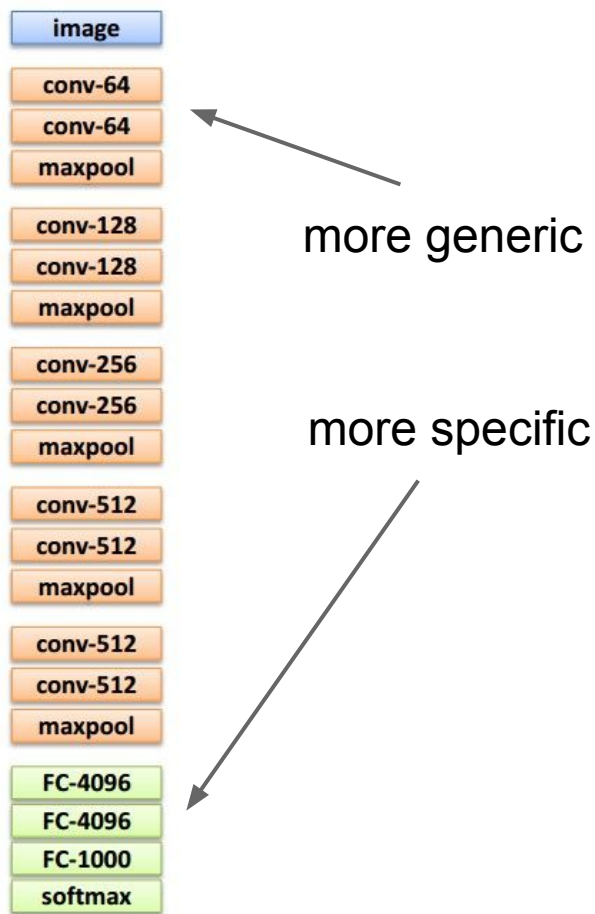


	very similar dataset	very different dataset
very little data	?	?
quite a lot of data	?	?





	<b>very similar dataset</b>	<b>very different dataset</b>
<b>very little data</b>	Use Linear Classifier on top layer	?
<b>quite a lot of data</b>	Finetune a few layers	?



	<b>very similar dataset</b>	<b>very different dataset</b>
<b>very little data</b>	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
<b>quite a lot of data</b>	Finetune a few layers	Finetune a larger number of layers

# Transfer learning with CNNs is pervasive... (it's the norm, not an exception)

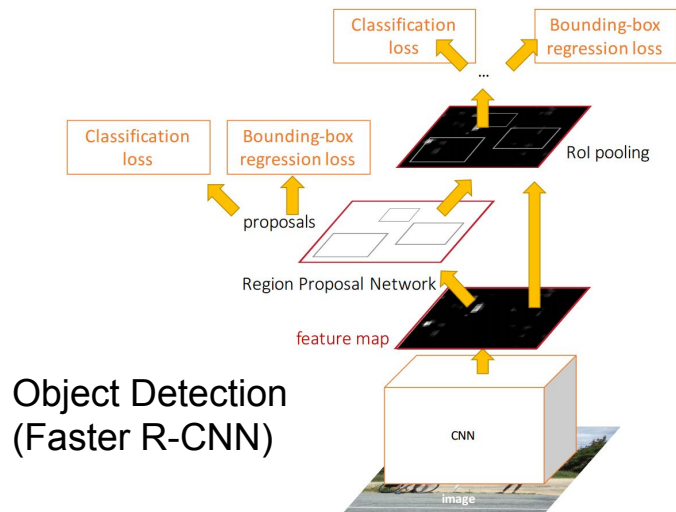

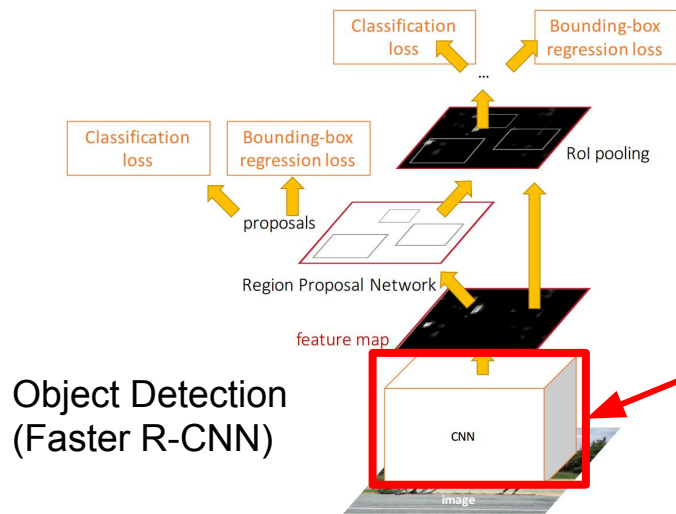


Image Captioning: CNN + RNN




# Transfer learning with CNNs is pervasive...

(it's the norm, not an exception)




CNN pretrained on ImageNet

Image Captioning: CNN + RNN



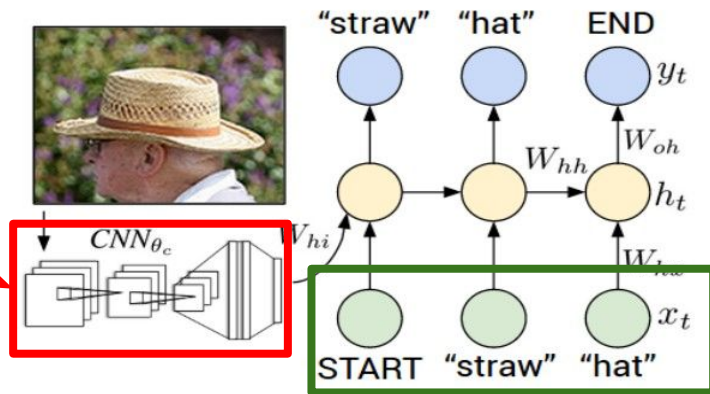
# Transfer learning with CNNs is pervasive...

(it's the norm, not an exception)



CNN pretrained  
on ImageNet

Image Captioning: CNN + RNN



Word vectors pretrained  
from word2vec

# Takeaway for your projects/beyond:

Have some dataset of interest but it has  $< \sim 1\text{M}$  images?

1. Find a very large dataset that has similar data, train a big ConvNet there.
2. Transfer learn to your dataset

Caffe ConvNet library has a “**Model Zoo**” of pretrained models:  
<https://github.com/BVLC/caffe/wiki/Model-Zoo>

# All About Convolutions


# All About Convolutions

## Part I: How to stack them




# The power of small filters

Suppose we stack two 3x3 conv layers (stride 1)  
Each neuron sees 3x3 region of previous activation map



# The power of small filters


**Question:** How big of a region in the input does a neuron on the second conv layer see?



# The power of small filters

**Question:** How big of a region in the input does a neuron on the second conv layer see?

**Answer:** 5 x 5




# The power of small filters

**Question:** If we stack **three** 3x3 conv layers, how big of an input region does a neuron in the third layer see?

# The power of small filters

**Question:** If we stack **three** 3x3 conv layers, how big of an input region does a neuron in the third layer see?


**Answer: 7 x 7**



# The power of small filters

**Question:** If we stack **three** 3x3 conv layers, how big of an input region does a neuron in the third layer see?

**Answer: 7 x 7**



Three 3 x 3 conv  
gives similar  
representational  
power as a single  
7 x 7 convolution

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

three CONV with  $3 \times 3$  filters

Number of weights:



# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:  
 $= C \times (7 \times 7 \times C) = \mathbf{49\ C^2}$

three CONV with  $3 \times 3$  filters

Number of weights:  
 $= 3 \times C \times (3 \times 3 \times C) = \mathbf{27\ C^2}$

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:  
 $= C \times (7 \times 7 \times C) = 49 C^2$

three CONV with  $3 \times 3$  filters

Number of weights:  
 $= 3 \times C \times (3 \times 3 \times C) = 27 C^2$



Fewer parameters, more nonlinearity = GOOD

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:  
 $= C \times (7 \times 7 \times C) = 49 C^2$

Number of multiply-adds:

three CONV with  $3 \times 3$  filters

Number of weights:  
 $= 3 \times C \times (3 \times 3 \times C) = 27 C^2$

Number of multiply-adds:

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:

$$= C \times (7 \times 7 \times C) = 49 C^2$$

Number of multiply-adds:

$$\begin{aligned} &= (H \times W \times C) \times (7 \times 7 \times C) \\ &= \mathbf{49 HWC^2} \end{aligned}$$

three CONV with  $3 \times 3$  filters

Number of weights:

$$= 3 \times C \times (3 \times 3 \times C) = 27 C^2$$

Number of multiply-adds:

$$\begin{aligned} &= 3 \times (H \times W \times C) \times (3 \times 3 \times C) \\ &= \mathbf{27 HWC^2} \end{aligned}$$

# The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:  
 $= C \times (7 \times 7 \times C) = 49 C^2$

Number of multiply-adds:  
 $= 49 HWC^2$

three CONV with  $3 \times 3$  filters

Number of weights:  
 $= 3 \times C \times (3 \times 3 \times C) = 27 C^2$

Number of multiply-adds:  
 $= 27 HWC^2$


 Less compute, more nonlinearity = GOOD

# The power of small filters

Why stop at  $3 \times 3$  filters? Why not try  $1 \times 1$ ?

# The power of small filters


Why stop at 3 x 3 filters? Why not try 1 x 1?



1. “bottleneck” 1 x 1 conv to reduce dimension

# The power of small filters

Why stop at 3 x 3 filters? Why not try 1 x 1?




1. “bottleneck” 1 x 1 conv to reduce dimension
2. 3 x 3 conv at reduced dimension



# The power of small filters

Why stop at 3 x 3 filters? Why not try 1 x 1?




1. “bottleneck” 1 x 1 conv to reduce dimension
2. 3 x 3 conv at reduced dimension
3. Restore dimension with another 1 x 1 conv

[Seen in Lin et al, “Network in Network”, GoogLeNet, ResNet]

# The power of small filters


Why stop at 3 x 3 filters? Why not try 1 x 1?



# The power of small filters

Why stop at 3 x 3 filters? Why not try 1 x 1?

More nonlinearity,  
fewer params,  
less compute!




# The power of small filters

Still using  $3 \times 3$  filters ... can we break it up?

# The power of small filters


Still using 3 x 3 filters ... can we break it up?



# The power of small filters

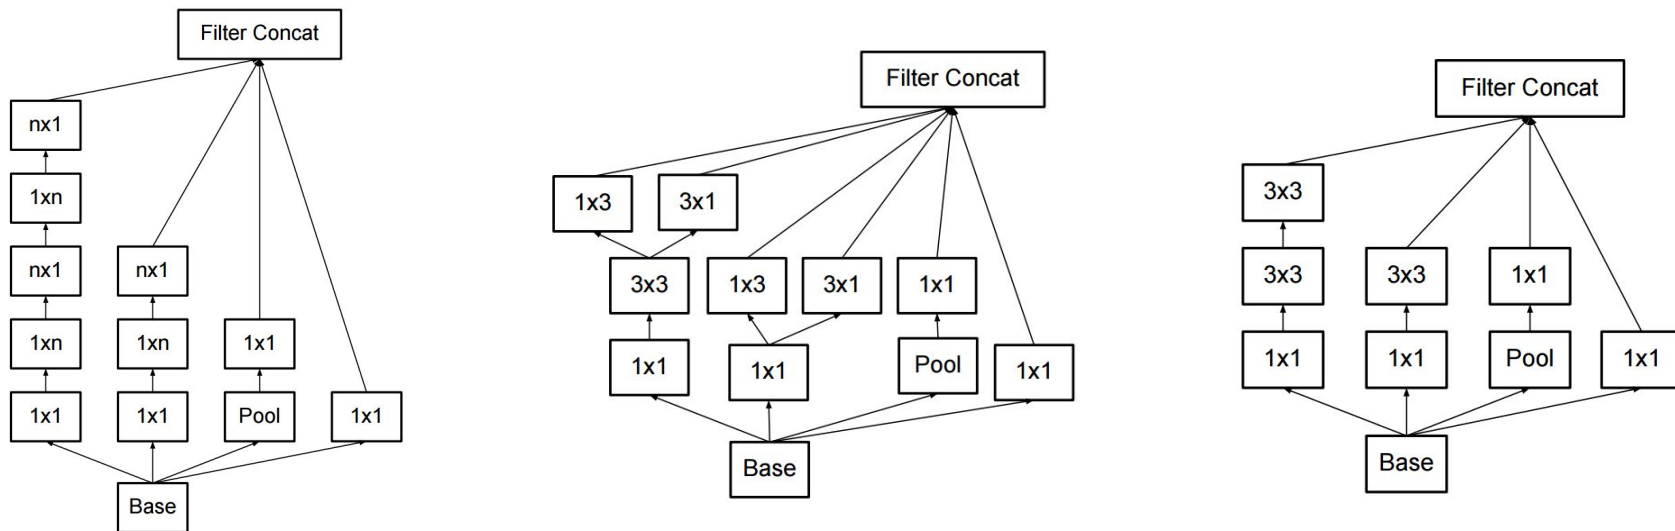
Still using 3 x 3 filters ... can we break it up?

More nonlinearity,  
fewer params,  
less compute!



# The power of small filters

Latest version of GoogLeNet incorporates all these ideas



Szegedy et al, "Rethinking the Inception Architecture for Computer Vision"

# How to stack convolutions: Recap

- Replace large convolutions ( $5 \times 5$ ,  $7 \times 7$ ) with stacks of  $3 \times 3$  convolutions
- $1 \times 1$  “bottleneck” convolutions are very efficient
- Can factor  $N \times N$  convolutions into  $1 \times N$  and  $N \times 1$
- All of the above give fewer parameters, less compute, more nonlinearity



# All About Convolutions

## Part II: How to compute them


# Implementing Convolutions: im2col

There are highly optimized matrix multiplication routines for just about every platform


Can we turn convolution into matrix multiplication?

# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$





Conv weights:  $D$  filters, each  $K \times K \times C$



# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$


Conv weights:  $D$  filters, each  $K \times K \times C$




# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$

Conv weights:  $D$  filters, each  $K \times K \times C$




Repeat for all columns to get  $(K^2C) \times N$  matrix  
( $N$  receptive field locations)




# Implementing Convolutions: im2col

Feature map:  $H \times W \times C$

Conv weights:  $D$  filters, each  $K \times K \times C$

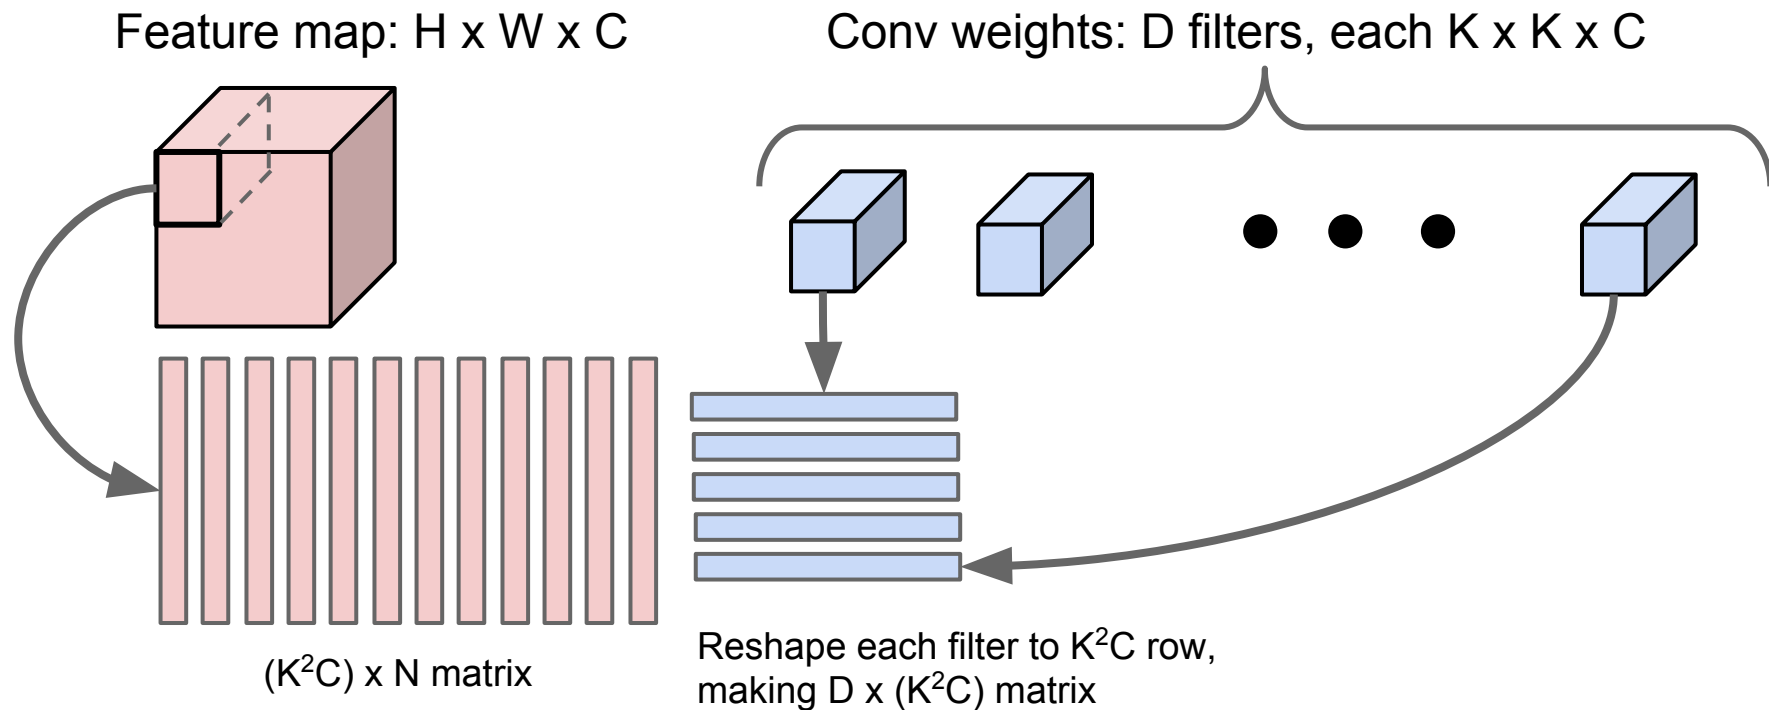


Repeat for all columns to get  $(K^2C) \times N$  matrix  
( $N$  receptive field locations)




Elements appearing in multiple  
receptive fields are duplicated; this  
uses a lot of memory

# Implementing Convolutions: im2col



# Implementing Convolutions: im2col





```

template <typename Dtype>
void ConvolutionLayer<Dtype>::Forward_gpu(const vector<Blob<Dtype>*>& bottom,
    vector<Blob<Dtype>*>* top) {
    for (int i = 0; i < bottom.size(); ++i) {
        const Dtype* bottom_data = bottom[i]->gpu_data();
        Dtype* top_data = (*top)[i]->mutable_gpu_data();
        Dtype* col_data = col_buffer_.mutable_gpu_data();
        const Dtype* weight = this->blobs_[0]->gpu_data();
        int weight_offset = M_ * K_;
        int col_offset = K_ * N_;
        int top_offset = M_ * N_;
        for (int n = 0; n < num_; ++n) {
            // im2col transformation: unroll input regions for filtering
            // into column matrix for multiplication.
            im2col_gpu(bottom_data + bottom[i]->offset(n), channels_, height_,
                width_, kernel_h_, kernel_w_, pad_h_, pad_w_, stride_h_, stride_w_,
                col_data);
            // Take inner products for groups.
            for (int g = 0; g < group; ++g) {
                caffe_gpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, M_, N_, K_,
                    (Dtype)1., weight + weight_offset * g, col_data + col_offset * g,
                    (Dtype)0., top_data + (*top)[i]->offset(n) + top_offset * g);
            }
            // Add bias.
            if (bias_term_) {
                caffe_gpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, num_output_,
                    N_, 1, (Dtype)1., this->blobs_[1]->gpu_data(),
                    bias_multiplier_.gpu_data(),
                    (Dtype)1., top_data + (*top)[i]->offset(n));
            }
        }
    }
}

```

## Case study: CONV forward in Caffe library

im2col

matrix multiply: call to  
cuBLAS

bias offset

```

def conv_forward_strides(x, w, b, conv_param):
    N, C, H, W = x.shape
    F, _, HH, WW = w.shape
    stride, pad = conv_param['stride'], conv_param['pad']

    # Check dimensions
    assert (W + 2 * pad - WW) % stride == 0, 'width does not work'
    assert (H + 2 * pad - HH) % stride == 0, 'height does not work'

    # Pad the input
    p = pad
    x_padded = np.pad(x, ((0, 0), (0, 0), (p, p), (p, p)), mode='constant')

    # Figure out output dimensions
    H += 2 * pad
    W += 2 * pad
    out_h = (H - HH) / stride + 1
    out_w = (W - WW) / stride + 1

    # Perform an im2col operation by picking clever strides
    shape = (C, HH, WW, N, out_h, out_w)
    strides = (H * W, W, 1, C * H * W, stride * W, stride)
    strides = x.itemsize * np.array(strides)
    x_stride = np.lib.stride_tricks.as_strided(x_padded,
                                                shape=shape, strides=strides)
    x_cols = np.ascontiguousarray(x_stride)
    x_cols.shape = (C * HH * WW, N * out_h * out_w)

    # Now all our convolutions are a big matrix multiply
    res = w.reshape(F, -1).dot(x_cols) + b.reshape(-1, 1)

    # Reshape the output
    res.shape = (F, N, out_h, out_w)
    out = res.transpose(1, 0, 2, 3)

    # Be nice and return a contiguous array
    # The old version of conv_forward_fast doesn't do this, so for a fair
    # comparison we won't either
    out = np.ascontiguousarray(out)

    cache = (x, w, b, conv_param, x_cols)
    return out, cache

```

## Case study: fast\_layers.py from HW

im2col

matrix multiply:  
call np.dot  
(which calls BLAS)

# Implementing convolutions: FFT

**Convolution Theorem:** The convolution of  $f$  and  $g$  is equal to the elementwise product of their Fourier Transforms:


$$\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$$

Using the **Fast Fourier Transform**, we can compute the Discrete Fourier transform of an  $N$ -dimensional vector in  $O(N \log N)$  time (also extends to 2D images)

# Implementing convolutions: FFT

1. Compute FFT of weights:  $F(W)$
2. Compute FFT of image:  $F(X)$
3. Compute elementwise product:  $F(W) \circ F(X)$
4. Compute inverse FFT:  $Y = F^{-1}(F(W) \circ F(X))$

# Implementing convolutions: FFT



FFT convolutions get a big speedup for larger filters  
Not much speedup for 3x3 filters =(

Vasilache et al, Fast Convolutional Nets With fbfft: A GPU Performance Evaluation

# Implementing convolution: “Fast Algorithms”

**Naive matrix multiplication:** Computing product of two  $N \times N$  matrices takes  $O(N^3)$  operations

**Strassen's Algorithm:** Use clever arithmetic to reduce complexity to  $O(N^{\log_2(7)}) \sim O(N^{2.81})$

$$\begin{array}{ll} A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} & \begin{array}{l} M_1 := (A_{1,1} + A_{2,2})(B_{1,1} + B_{2,2}) \\ M_2 := (A_{2,1} + A_{2,2})B_{1,1} \\ M_3 := A_{1,1}(B_{1,2} - B_{2,2}) \\ M_4 := A_{2,2}(B_{2,1} - B_{1,1}) \\ M_5 := (A_{1,1} + A_{1,2})B_{2,2} \\ M_6 := (A_{2,1} - A_{1,1})(B_{1,1} + B_{1,2}) \\ M_7 := (A_{1,2} - A_{2,2})(B_{2,1} + B_{2,2}) \end{array} \\ B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} & \\ C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix} & \begin{array}{l} C_{1,1} = M_1 + M_4 - M_5 + M_7 \\ C_{1,2} = M_3 + M_5 \\ C_{2,1} = M_2 + M_4 \\ C_{2,2} = M_1 - M_2 + M_3 + M_6 \end{array} \end{array}$$

From Wikipedia

# Implementing convolution: “Fast Algorithms”

Similar cleverness can be applied to convolutions

Lavin and Gray (2015) work out special cases for 3x3 convolutions:

$$F(2,3) = \begin{bmatrix} d_0 & d_1 & d_2 \\ d_1 & d_2 & d_3 \end{bmatrix} \begin{bmatrix} g_0 \\ g_1 \\ g_2 \end{bmatrix} = \begin{bmatrix} m_1 + m_2 + m_3 \\ m_2 - m_3 - m_4 \end{bmatrix}$$

$$\begin{aligned} m_1 &= (d_0 - d_2)g_0 & m_2 &= (d_1 + d_2)\frac{g_0 + g_1 + g_2}{2} \\ m_4 &= (d_1 - d_3)g_2 & m_3 &= (d_2 - d_1)\frac{g_0 - g_1 + g_2}{2} \end{aligned}$$

$$\begin{aligned} B^T &= \begin{bmatrix} 1 & 0 & -1 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & 1 & 0 & -1 \end{bmatrix} \\ G &= \begin{bmatrix} 1 & 0 & 0 \\ \frac{1}{2} & \frac{1}{2} & \frac{1}{2} \\ \frac{1}{2} & -\frac{1}{2} & \frac{1}{2} \\ 0 & 0 & 1 \end{bmatrix} \\ A^T &= \begin{bmatrix} 1 & 1 & 1 & 0 \\ 0 & 1 & -1 & -1 \end{bmatrix} \\ g &= [g_0 \ g_1 \ g_2]^T \\ d &= [d_0 \ d_1 \ d_2 \ d_3]^T \end{aligned}$$

Lavin and Gray, “Fast Algorithms for Convolutional Neural Networks”, 2015

# Implementing convolution: “Fast Algorithms”

Huge speedups on VGG for small batches:

N	cuDNN		F(2x2,3x3)		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	12.52	3.12	5.55	7.03	2.26X
2	20.36	3.83	9.89	7.89	2.06X
4	104.70	1.49	17.72	8.81	5.91X
8	241.21	1.29	33.11	9.43	7.28X
16	203.09	3.07	65.79	9.49	3.09X
32	237.05	5.27	132.36	9.43	1.79X
64	394.05	6.34	266.48	9.37	1.48X

Table 5. cuDNN versus  $F(2 \times 2, 3 \times 3)$  performance on VGG Network E with fp32 data. Throughput is measured in Effective TFLOPS, the ratio of direct algorithm GFLOPs to run time.

N	cuDNN		F(2x2,3x3)		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	14.58	2.68	5.53	7.06	2.64X
2	20.94	3.73	9.83	7.94	2.13X
4	104.19	1.50	17.50	8.92	5.95X
8	241.87	1.29	32.61	9.57	7.42X
16	204.01	3.06	62.93	9.92	3.24X
32	236.13	5.29	123.12	10.14	1.92X
64	395.93	6.31	242.98	10.28	1.63X

Table 6. cuDNN versus  $F(2 \times 2, 3 \times 3)$  performance on VGG Network E with fp16 data.



# Computing Convolutions: Recap

- im2col: Easy to implement, but big memory overhead
- FFT: Big speedups for small kernels
- “Fast Algorithms” seem promising, not widely used yet

# Implementation Details



# Spot the CPU!





# Spot the CPU!

“central processing unit”



# Spot the GPU!

“graphics processing unit”






# Spot the GPU!

“graphics processing unit”






vs







VS

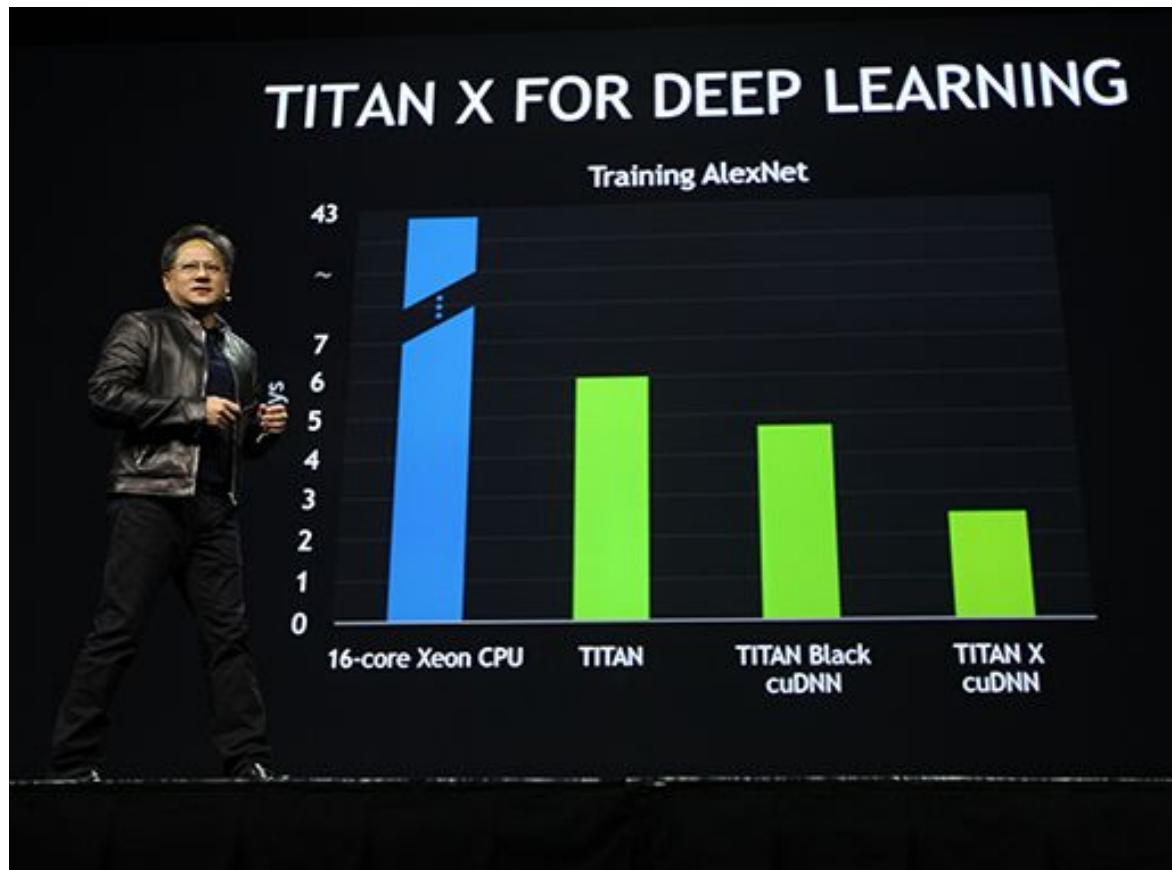


NVIDIA is much more  
common for deep learning

**CEO of NVIDIA:**  
Jen-Hsun Huang

(Stanford EE Masters  
1992)

**GTC 2015:**  
Introduced new Titan X  
GPU by bragging about  
AlexNet benchmarks



## CPU

Few, fast cores (1 - 16)

Good at sequential processing



## GPU

Many, slower cores (thousands)

Originally for graphics


Good at parallel computation



# GPUs can be programmed

- CUDA (NVIDIA only)
  - Write C code that runs directly on the GPU
  - Higher-level APIs: cuBLAS, cuFFT, cuDNN, etc
- OpenCL
  - Similar to CUDA, but runs on anything
  - Usually slower :(
- Udacity: Intro to Parallel Programming <https://www.udacity.com/course/cs344>
  - For deep learning just use existing libraries


GPUs are really good  
at matrix multiplication:



← **GPU:** NVIDIA Tesla K40  
with cuBLAS

← **CPU:** Intel E5-2697 v2  
12 core @ 2.7 Ghz  
with MKL

## GPUs are really good at convolution (cuDNN):




All comparisons are against a 12-core Intel E5-2679v2 CPU @ 2.4GHz running Caffe with Intel MKL 11.1.3.

Even with GPUs, training can be slow

**VGG:** ~2-3 weeks training with 4 GPUs

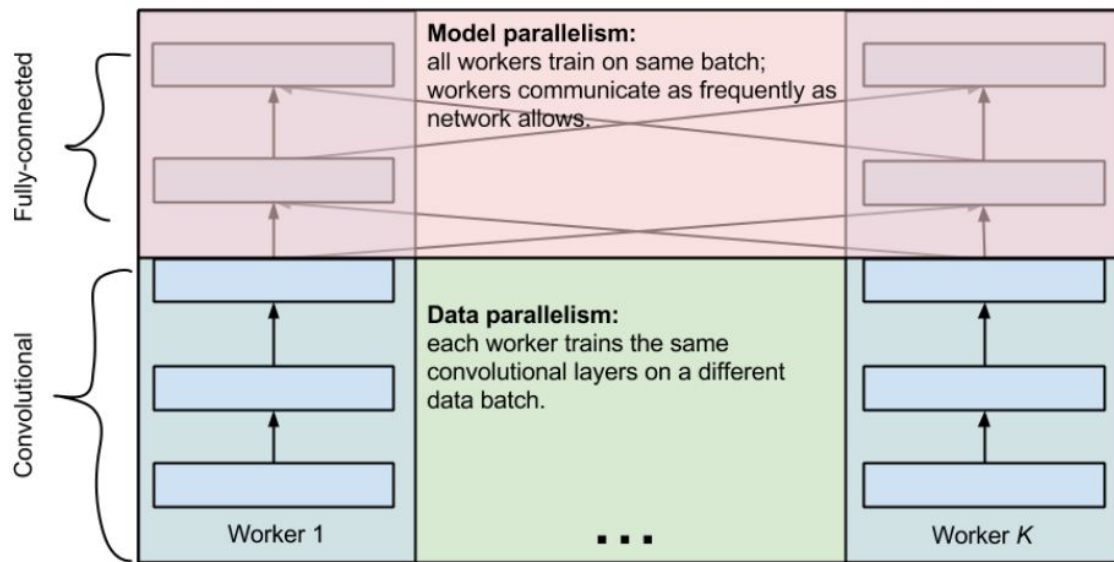
**ResNet 101:** 2-3 weeks with 4 GPUs

NVIDIA Titan Blacks  
~\$1K each



ResNet reimplemented in Torch: <http://torch.ch/blog/2016/02/04/resnets.html>


# Multi-GPU training: More complex



Alex Krizhevsky, "One weird trick for parallelizing convolutional neural networks"




# Google: Distributed CPU training




**Data parallelism**

*[Large Scale Distributed Deep Networks, Jeff Dean et al., 2013]*

# Google: Distributed CPU training



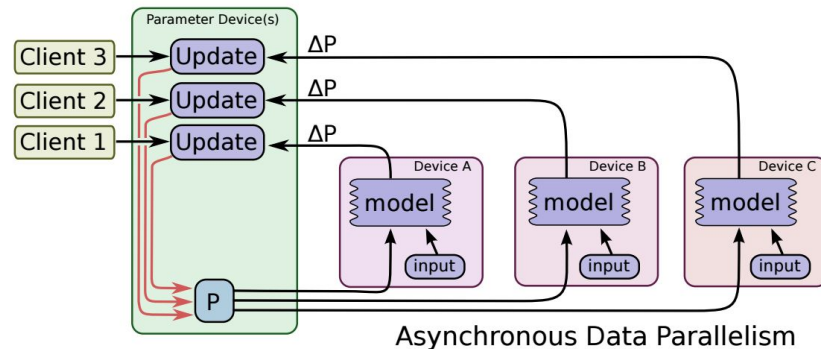
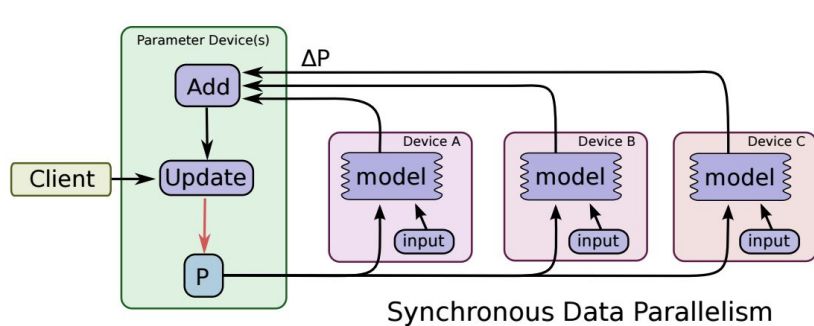
**Data parallelism**



**Model parallelism**

*[Large Scale Distributed Deep Networks, Jeff Dean et al., 2013]*

# Google: Synchronous vs Async



*Abadi et al, "TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems"*

# Bottlenecks

to be aware of



**GPU - CPU communication is a bottleneck.**

=>

**CPU** data prefetch+augment thread running

while

**GPU** performs forward/backward pass

# CPU - disk bottleneck

Hard disk is slow to read from

=> Pre-processed images  
stored contiguously in files, read as  
raw byte stream from SSD disk

Moving parts lol



# GPU memory bottleneck

Titan X: 12 GB <- currently the max  
GTX 980 Ti: 6 GB

e.g.

AlexNet: ~3GB needed with batch size 256

# Floating Point Precision



# Floating point precision

- 64 bit “double” precision is default in a lot of programming
- 32 bit “single” precision is typically used for CNNs for performance

# Floating point precision

- 64 bit “double” precision is default in a lot of programming
- 32 bit “single” precision is typically used for CNNs for performance
  - Including cs231n homework!

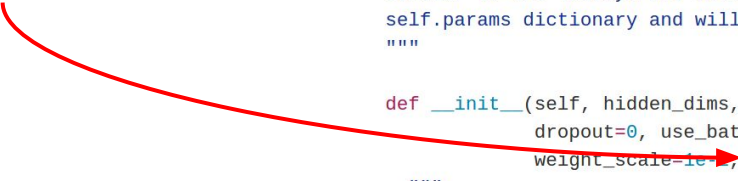
```
class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=0, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-4, dtype=np.float32, seed=None):
```



# Floating point precision

Benchmarks on Titan X, from <https://github.com/soumith/convnet-benchmarks>

**Prediction:** 16 bit “half” precision will be the new standard

- Already supported in cuDNN
- Nervana fp16 kernels are the fastest right now
- Hardware support in next-gen NVIDIA cards (Pascal)
- Not yet supported in torch =(

AlexNet (One Weird Trick paper) - Input 128x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	92	29	62
CuDNN[R3]-fp16 (Torch)	<a href="#">cudnn.SpatialConvolution</a>	96	30	66
CuDNN[R3]-fp32 (Torch)	<a href="#">cudnn.SpatialConvolution</a>	96	32	64

OxfordNet [Model-A] - Input 64x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	529	167	362
Nervana-fp32	ConvLayer	590	180	410
CuDNN[R3]-fp16 (Torch)	<a href="#">cudnn.SpatialConvolution</a>	615	179	436

GoogleNet V1 - Input 128x3x224x224



Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	283	85	197
Nervana-fp32	ConvLayer	322	90	232
CuDNN[R3]-fp32 (Torch)	<a href="#">cudnn.SpatialConvolution</a>	431	117	313

# Floating point precision

How low can we go?

Gupta et al, 2015:

Train with **16-bit fixed point** with stochastic rounding



CNNs on MNIST

Gupta et al, "Deep Learning with Limited Numerical Precision", ICML 2015

# Floating point precision

How low can we go?

Courbariaux et al, 2015:

Train with **10-bit activations**, **12-bit parameter updates**

Courbariaux et al, "Training Deep Neural Networks with Low Precision Multiplications", ICLR 2015

# Floating point precision

How low can we go?

Courbariaux and Bengio, February 9 2016:

- Train with **1-bit activations and weights!**
- All activations and weights are +1 or -1
- Fast multiplication with bitwise XNOR
- (Gradients use higher precision)

Courbariaux et al, “BinaryNet: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1”, arXiv 2016

# Implementation details: Recap

- GPUs much faster than CPUs
- Distributed training is sometimes used
  - Not needed for small problems
- Be aware of bottlenecks: CPU / GPU, CPU / disk
- Low precision makes things faster and still works
  - 32 bit is standard now, 16 bit soon
  - In the future: binary nets?

# Recap

- Data augmentation: artificially expand your data
- Transfer learning: CNNs without huge data
- All about convolutions
- Implementation details